

A Collaborative guide to ATLAS Development

R. Clint Whaley ^{*}Peter Soendergaard [†]

August 4, 2014

Abstract

This paper exists to get you started if you want to do some ATLAS development. The hope is that as new developers work on ATLAS, they will add to this note, so that this information grows with time.

^{*}`rwhaley@users.sourceforge.net`

[†]`soender@users.sourceforge.net`

Contents

1	Introduction	3
2	Adding to this note	3
3	Getting the ATLAS code through github	4
3.1	Background on ATLAS code base	4
3.2	Getting the ATLAS basefile tree and the working ATLAS directory	4
3.3	Basefile/extracted file interaction	5
4	ATLAS coding style	5
5	Adding a LAPACK routine to ATLAS	7
5.1	Row-major LAPACK routines	7
5.2	Outline of Steps	8
5.3	Create and debug tester	8
5.3.1	Writing ATLAS/src/testing C-to-f77 wrapper	9
5.3.2	Getting your routines extracted	9
5.4	Create and debug ATLAS internal routines	9
5.5	Add C and F77 interface routines	10
5.6	Update the LAPACK quick reference guides	10
6	Architectural defaults	10
6.1	Rambling on about architectural defaults	10
6.2	Making your own architectural defaults	11
6.3	Getting ATLAS to use your shiny new defaults	11
7	Sanity testing for an ATLAS install	12
7.1	Invoking the sanity tests	12
7.2	Understanding the sanity test output	12
7.3	Finding the context of the error	13
7.4	Tracking down an error in the BLAS interface testers	14
7.5	Tracking down an error in the bin/ testers	15
8	Antoine's testing scripts	15
8.1	Finding errors	16
8.2	Tracking down errors in the bin/ testers	16
9	Finding a good NB for GEMM	16
9.1	Better kernel timing does not always yield faster GEMM	17
9.2	Large NB means more time in cleanup	17
9.3	Large NB means more time in unblocked application code	17
9.4	Finding a good NB	18

10 Understanding the ATLAS configure system	20
10.1 Probe Overview	20
10.2 Architectural Probes	22
10.3 Notes on configure	23
10.4 Compiler Setup and Handling in ATLAS Config	23
10.5 Notes on cross-compilation:	24
10.6 Notes on native Windows compilation:	24

1 Introduction

This note exists to get you started if you want to write **new** code for ATLAS, or if you want to modify ATLAS source. It is *not* for kernel contribution, which is what most people do when contributing to ATLAS. Kernel contribution is much simpler, and is explained in `ATLAS/doc/atlas_contrib.pdf` or in html format at:

http://math-atlas.sourceforge.net/devel/atlas_contrib/

So what is the difference between development and kernel contribution? In kernel contribution, you write a kernel to be used by ATLAS, using the provided ATLAS testers and timers to verify it, and when you are satisfied with its performance and reliability, you submit your kernel to the ATLAS team, and they accept it or not, and you are done.

Doing actual core development is quite a bit more complex. Probably the biggest change is that you will need to write your own tester and timer for your new code. No code will be accepted into the ATLAS code base without a tester which can be used to verify it. Since writing a decent tester is usually at least as hard as writing the code it tests, and is always a whole lot less enjoyable, the author must bear the pain of producing it along with the pride of producing the code. As a developer, you will be responsible for testing such new code on several platforms as well.

If you are instead hoping to modify some of the existing code base, remember that for non-kernel operations, portability and robustness *must* be the primary goal. There are many sections of ATLAS that we know to be second rate on a certain platform, but we also know that it works on the twenty or so architectures that ATLAS is routinely compiled for, so we leave it that way. This means that when a modification is made to a previously existing routine, the modifying author must have good evidence that the new code is as portable as the old. In short, the barrier to replacing tested code is high.

It is possible that users want access to the git repository even though they do not plan on doing development, mainly I'd guess so they can get access to the newest stuff without waiting for developer releases. Also, kernel contributors who make subsequent changes to their routines can speed up their adoption by submitting them as patches on the basefiles, or through a pull request.

2 Adding to this note

This note is included in the `AtlasBase/TeXDoc` of the repository, and anyone can submit a patch against it giving additional information. As the founder of ATLAS, I have written a seed of a document explaining how to get access to the code. It is my hope that other developers will add important information that they discover as they go, so that this doc will grow over time, getting much information that I probably take so much for granted that I would never think to document.

Therefore, new sections are welcome, and probably a FAQ appendix would be a good idea. As people contribute, their names will be added to the author list.

3 Getting the ATLAS code through github

3.1 Background on ATLAS code base

ATLAS was not originally developed in with version control. ATLAS was developed using a programming tool called *extract*, which means ATLAS is actually maintained in something called basefiles. If you think of regular development being access by value, and version control using a level of indirection, version control on basefiles gives you two such levels of indirection. So, if you want to be able to submit directly to the repository, you will need to learn at least the basics of *extract*. Details on *extract* can be found at:

<http://www.cs.utsa.edu/~whaley/extract/Extract.html>

Note that if you just want read-only access, the install process will use *extract*, but you will not need to learn anything about it.

3.2 Getting the ATLAS basefile tree and the working ATLAS directory

Here are the steps to get a local ATLAS basefile tree:

1. If you don't already have it, install git. On my lubuntu system, I did: “`sudo apt-get install -y git-core gitg qgit`”.
2. **Checkout the ATLAS basefiles:** From a git-capable machine connected to the Internet, go to the directory where you want the ATLAS basefiles to be, and issue:

```
git clone git://github.com/rcwhaley/math-atlas.git
```

This command should create a directory called `math-atlas` in the current directory, and `math-atlas` should contain the files `Makefile`, `extract.c`, and `README`, as well as the directory `AtlasBase`, which contains all the ATLAS basefiles, and the directory `AtlasTest`, which contains Antoine's testing scripts.

→ The details of how to use git actually vary depending on what you are trying to do, and the above command is just the simplest thing. Please see:

<http://nipy.sourceforge.net/nipy/stable/devel/guidelines/gitwash/>
for a good discussion of the two primary ways to use git to contribute: (1) submitting patches, or (2) submitting substantial new development with forking and pull requests.

3. **Create ATLAS working directory and extract files:** Decide where you wish to create the ATLAS source tree. This directory should not presently exist, since the `Makefile` is going to create it for you. Create the source tree with the command:

```
make srcdir=<your chosen directory>
```

If `srcdir` is not given, the `Makefile` creates a directory `TEST` in the `math-atlas` directory (i.e., `srcdir` defaults to `./TEST`).

Now, `cd` to `srcdir` and issue ‘`make`’ to build the ATLAS tree, and `./atltar.sh` will build the standard ATLAS tarfile from the ATLAS source tree.

3.3 Basefile/extracted file interaction

You'll have to scope the extract page for any kind of real feel for how this works, but some atlas-specific details are in order here. In each subdirectory of the `ATLAS/` tree, you will find a file `Make.ext`. If you type `make -f Make.ext`, this makefile will extract all new files in this subdirectory for which the basefile is newer.

So, what usually happens is you are messing with something, and you do it in the `ATLAS/` directory. When you are confident in your change, you put it into the appropriate basefile in the `AtlasBase/` directory (note that examining `Make.ext` will show you what basefile a given extracted file comes from), and you then re-extract over your working copy with the above command.

4 ATLAS coding style

ATLAS attempts to use a unified coding style. Older codes may not adhere to it strictly (especially in function prototyping and single-line ifs), but all new codes should do so.

- Strict ANSI/ISO 9899-1990 C conformance is required
 - Must use `/* */`, not `//`
- Lines are at most 80 characters in length
- No tabs should be used in the file, only spaces
 - in vim can use `:set expandtab`
- Indentation is always 3 characters, and the braces should line up, as in:

```
for (i=0; i < N; i++)
{
    for (j=0; j < N; j++)
        statement;
}
```

All braces are lined up this way (eg., ifs, functions & loops).

- Multiline function calls have arguments indented to the opening paren:

```
error = function_call(arg1, arg2, arg3,
                      arg4, arg5);
```

- If bodies are always on a separate line from if, so you want:

```
if (cond)
    statement;
```

rather than:

```
if (cond) statement;
```

- Comments are either to the right of a line if there is room, or on lines preceeding the commented code if not, using a style like:

```
/*
 *      This comment describes what is going on the below loop body
 *      and if conditional
 */
    if (cond)
    {
    }

/*
 *      Comment describing else body
 */
    else
    {
    }
```

Note that the comment text is lined up with the code it comments.

- Occasionally you need a superblock of comments, that describes a whole region of code (eg, multiple if/else blocks). This is done as:

```
/*
 * =====
 * Here is a comment describing the total operation of multiple blocks of
 * code which are all individually commented in the above manner
 * =====
 */

    for ()
    {
        if (cond)
            ....
        else if
            ...
    }
```

- Functions prologues are commented like this:

```
int my_func
(
    int iparm,          /* comment describing this parameter */
    void *vp,           /* comment describing this paraemter */
    ...
)
/*
 * This mandatory block of comments placed between the func declaration and
 * the opening brace describes what the function does.  If the function
```

```

* returns a value, this block of comments end with line(s) describing the
* the return value, of the form:
* RETURNS: 0 on success, non-zero error message otherwise
*         second line describing return value
*/
{
}

```

5 Adding a LAPACK routine to ATLAS

It is probably not practical that ATLAS will provide a complete LAPACK API (as it does with the BLAS) in the foreseeable future, both due to the algorithmic complexity of some of the operations, and to the sheer number of routines in LAPACK. It must be understood that adding routines adds to the inertia and maintenance costs of the package, and this additional burden must be offset by real advantage for the user.

ATLAS has so far only added LAPACK routines to ATLAS when we can make a performance-enhancing algorithm modification. For instance, we added the LU and Cholesky factorizations because we used the recursive formulations of these routines, which provides for better performance on pretty much any cache-based architecture.

However, when we have added such routines, we usually add the correlated routines even when a performance advantage is not supplied. For instance, upon adding GETRF support, we also added GETRS and GESV. As far as column-major routines go, we supply no better algorithm for GETRS or GESV than LAPACK. However, since these routines are very simple, and GETRF is very often used with them, we added them along with GETRF. The idea here is that their maintenance costs are not heavy, and real advantage is given to the user in that we have sped up GETRF, and if the factor and solve are all he needs, ATLAS will supply a complete solution.

The column-major comment points out another reason to add a routine to ATLAS: ATLAS supplies the only performance-aware row-major LAPACK implementation that I am aware of (I'm sure there are some, I just don't know of any that aren't simply using the col-major stuff, and thus performing terribly). It is possible that someone would want to add an LAPACK routine to ATLAS simply because they need a row-major version, and someone being motivated enough to write it would probably be ample justification to add the routine to the ATLAS tarfile.

5.1 Row-major LAPACK routines

So far, we have accepted no routines that do not also include a row-major equivalent, both for BLAS and for LAPACK. We hope to continue this. There are as yet only a few users of the row-major LAPACK/BLAS that I am aware of, but I believe that this is a chicken/egg problem.

Some people insist on using row-major arrays in C, but if they have access to a BLAS/LAPACK that supports it, they find the performance is no better than what they get with simple loops, or that it is calling the col-major in a naive way, and cutting the problems size they can solve in half by copying. Therefore, people with row-major bias don't call the stuff 'cause it doesn't help them, and the problem continues.

It is my belief, therefore, that good-quality row-major stuff must be produced before significant demand will appear. If I'm wrong, I guess we'll someday drop support for row-major, but I don't think this will be the case over a long enough time line.

Therefore, despite it being a hassle, having a good quality row-major implementation is critical for getting an LAPACK routine into ATLAS. For many routines, since we have row-major BLAS, the algorithm stays the same, and only some pointer arithmetic need be changed.

Other routines in LAPACK (GETRF is one) have a built in algorithmic bias towards column-major (in GETRF, this is doing row-pivoting), and another algorithm with the same stability and usage characteristics should be employed for row-major (eg., column-pivoting, for GETRF).

5.2 Outline of Steps

Here are the general steps to use when adding an LAPACK routine to ATLAS:

1. Create and debug tester using LAPACK
 - Update extractors
 - Update Makefiles
2. Write and test ATLAS internal routines using above tester
 - Update extractors
 - Update Makefiles
3. Update `atlas_clapack.h` to include your new routines
4. Create C and F77 interfaces to your routine
 - Update extractors
 - Update Makefiles
5. Update `clapack.h`
6. Update the LAPACK quick reference guides.

5.3 Create and debug tester

The first step in adding a new routine to ATLAS is to create a tester (and timer) which can be used to verify the correctness of your code. More than half of the challenge is getting the tester right; with a good tester/timer, the code usually comes fairly easily.

Your tester will go in `ATLAS/bin` when extracted; you can examine some of the testers available there to get an idea of what you should do (eg., look at `ATLAS/bin/[lu/llt/slv/trtri/uum]tst.c`). All of these routines come from the basefile `AtlasBase/Clint/atlas-tlp.base`, which is what you should submit your patch against, unless you want to create your own, separate basefile.

After your tester is written, its column-major components can be tested against LAPACK by using the `make <rou>tstF` target in `ATLAS/bin/<arch>`. You can even test the

row-major components by having the F77 interface transpose the matrices on input, and back on output. See `ATLAS/bin/uumtst.c` for an example of this for square matrices.

As part of your debugging of the tester, be sure that it not only agrees that LAPACK produces the right answer, but truly detects errors as well. For instance, manually overwrite an entry, both in the matrix and in the padding (in separate tests), and make sure it is caught by the tester.

5.3.1 Writing ATLAS/src/testing C-to-f77 wrapper

You first need a way for your tester, written in C, to call the LAPACK routine, written in F77. All such language translation routines are kept in `ATLAS/src/testing`, and come from the basefile `ATLAS/Clint/atlas-ilp.base`. This wrapper is trivial, though some of the integer/string stuff is not obvious. Steal the code from the other examples.

5.3.2 Getting your routines extracted

Now you need to get your files to appear in the right subdirectories, so you need an entry in the appropriate `Make.ext`. All the `Make.ext` files comes from `AtlasBase/make.base`, so find the rout for your directory in this file (for examples the line saying `@ROUT ATLAS/interfaces/lapack/F77/src/`) and add your routine name to the line containing the name of all the other routines.

So, now you are in your working directory (say `ATLAS/src/lapack`), and you type `make -f Make.ext`, and nothing happens, no new files show up. This is because you need to re-extract your `Make.ext` file. This can of course be done by removing your whole ATLAS tree and reinstalling, but less brutally you can “just” use something like this: `extract -b /home/soender/AtlasBase/make.base -o Make.ext rout=ATLAS/src/lapack -langM`. The `-langM` switch is required for `extract` to properly handle makefiles, so you cannot skip it.

This is the basic procedure for this sort of stuff. When you need a makefile in a `BLDdir` subdirectory, the appropriate makefile is copied by `Make.top` from the `ATLAS/makes/` directory. Check `Make.ext` to see which basefile they come from, and add your routine name among the names of the other routines.

Remember to update the Makefiles for both `ATLAS/bin` and `ATLAS/src/testing`, and to get these makefiles into the appropriate subdirs. In order to extract new makefiles, and get them put into the appropriate subdirs, I typically do something like (from the `BLDdir`:

```
pushd ~/TEST/ATLAS/makes/ ; make -f Make.ext ; cd .. ; \
    make refresh ; popd
```

(replace the path and arch appropriately, obviously).

5.4 Create and debug ATLAS internal routines

The internal LAPACK routines are kept in `AtlasBase/Clint/atlas-lp.base`. Add your routine here, and update `ATLAS/src/lapack's` `Make.ext` and `Makefile` appropriately to build your routine.

You will add your routine in `atlas-lp.base` with an additional `@ROUT` keyline, but also do not forget to update the include file `atlas_lapack.h` at the bottom of the file as well.

You will need to add your routine to the prototype part, as well as to the macro renaming part. Examine the basefile for details.

Once it is extracting, use your LAPACK-debugged tester to debug your code.

5.5 Add C and F77 interface routines

We do this step last, because we don't want to add API routines until the code is working. Having debugged and made sure the code is faster than LAPACK, we're now ready to make it available to the user via the advertised APIs. The extracted API files are kept in subdirectories under `ATLAS/interfaces/lapack`.

The F77 interfaces are kept in `AtlasBase/Clint/atlas-fint.base`. Look at the existing examples and notice how `extract` generates all four precision from the same routine, if you use the `extract` macros. All the code for this interface can be ripped from LAPACK and adapted. Note that you will usually need to examine both complex and real versions of the original LAPACK routine, to find any differences in interface/testing and comments. You will also need to remove unneeded `EXTERNAL` declaration, etc.

This interface does the parameter checking, and converts any FORTRAN string arguments to some predefined integer values, and then call the `ATLf77wrap` interface. Scope any of the existing routines for details on this.

The C interfaces are easy to write, since they should just check the input arguments, and then call the ATLAS routine. The codes are stored in `atlas-clp.base`. Check it out for lots of examples.

5.6 Update the LAPACK quick reference guides

The ATLAS user API is defined in the quick reference guides under `AtlasBase/TexDoc`. Right now, the supported LAPACK API is small enough to fit both C and F77 interfaces on one card (single 2-sided landscape page), but eventually it will be split in two, as with the BLAS quick reference cards. Either use the Makefile to do it, or remember to manually throw the `-tlandscape` flag to `dvips`, and the `-paper a4r` flag to `xdvi`.

6 Architectural defaults

ATLAS's architectural defaults are simply a record of the results of a previously run ATLAS search. They exist for a couple of reasons:

1. Using architectural defaults, install times are reduced to almost bearable levels
2. Because the search is empirical, installs can go wrong if unmonitored. Architectural defaults given out in the standard tarfile have at least passed the laugh test

6.1 Rambling on about architectural defaults

One FAQ for architectural defaults is why any timings are necessary when using them. The standard architectural defaults only rarely describe everything discovered by a search, but rather give only those data that we feel sure will not vary a great deal. For instance, for many machines, the kernels to use, etc., are fully specified, but `CacheEdge` is not. `CacheEdge`

varies depending on your L2 cache size, which varies depending on architecture revision, so it is not specified, allowing it to tune itself for this variable parameter, while still skipping the search over less variable things (eg., if the L1 cache or FPU units change, this is usually a new architecture, not a revision of an old).

That's the theoretical reason why they shouldn't cover all discovered items. However, ATLAS presently times the kernels in order to be able to produce a comprehensive SUMMARY.LOG, and these timings *could* be skipped, assuming this functionality were added to the atlas install process.

There are some weaknesses of architectural defaults. One of the main ones is how they can go out of date, and cause slowdown. One big way this can happen is with compiler changes. For instance, gcc 3.0 produces completely different (and inferior) x86 code than the 2.x series, and 4.0 was similarly worse than latter-day gcc 3. Almost all architectural defaults in ATLAS 3.10 were compiled with gcc 4.7.0.

Anytime a different compiler is used, the architectural defaults become suspect. For truly inferior compiler (like gcc 3.0, 4.0, or 4.1), there is no way to get good performance, but at least some problems can be worked around by having ATLAS adapt itself to the new compiler, and architectural defaults prevent this from happening.

6.2 Making your own architectural defaults

This section describes how to create architectural defaults as of ATLAS 3.9.80 and later. For older releases, the process is similar, but not quite the same, and is covered in the older `atlas_devel` available in those tarfiles.

1. Get an install, correct in all details, that you want to immortalize.
2. `cd` to your `OBJdir/ARCHS` directory
3. Type `make ArchNew`

This will copy the search result output files into a directory `<OBJdir>/ARCHS/<MACH>/`, with appropriate subdirs under that. You can then go into these guys and delete files you don't want to be part of the defaults (eg., `atlas_cacheedge.h`, etc).

Now, to save these defaults to a transportable format, you can have the makefile create the tarfile for you by:

```
make tarfile
```

6.3 Getting ATLAS to use your shiny new defaults

If necessary, rename the tarfile with your archdefs to the same name as the architecture you are presently compiling on (normally the file will already have the same name). Then, during configure, you simply tell ATLAS to use your directory of architectural defaults, by adding the following flag to configure:

```
-Ss AAddir <path to directory containing your tarfile>
```

Note that with this flag and ability to make archdefs, you can create an entire suite of custom archdefs for your particular machines, and keep them all together in a directory for future installs.

7 Sanity testing for an ATLAS install

From ATLAS3.3.8 forward, ATLAS has had a “sanity test”, which just does some quick testing in order to ensure that there are no obvious problems with the installed ATLAS libraries. It runs all of the standard BLAS interface testers, with the default input files, and it then runs a few fixed cases of ATLAS’s lapack tester routines (eg., `ATLAS/bin/invstst.c`, etc). The advantage of these lapack testers is that they depend on many of the BLAS as well as the lapack routines, so you get a lot of testing for a minor amount of time. The sanity checks do not require any non-ATLAS libraries for testing, so the only dependence that a user who has installed ATLAS may not be able to satisfy is the need for a Fortran77 compiler, which is required for the BLAS interface testers. As of ATLAS3.7.12, ATLAS can also run a reduced set of tests for users who do not have a fortran compiler.

7.1 Invoking the sanity tests

These tests are invoked from your install directory by:

```
make check
```

If you are using threads, you will want to run the same tests for threading via:

```
make ptcheck
```

7.2 Understanding the sanity test output

Once you fire off this tester, you’ll see a lot of compilation going on. All compilation is done up front, and then the testers are run at the end. All tester output is dumped to some files (we’ll see specifics in a bit), which are then automatically grepped for errors at the end of the run. It is the results of this grep that the user will see. For example, here’s the output from a run on my Athlon running Linux:

```
dudley.home.net. make check
...
... bunch of compilation ...
...
DONE BUILDING TESTERS, RUNNING:
SCOPING FOR FAILURES IN BIN TESTS:
fgrep -e fault -e FAULT -e error -e ERROR -e fail -e FAIL \
    bin/Linux_ATHLON/sanity.out
8 cases: 8 passed, 0 skipped, 0 failed
4 cases: 4 passed, 0 skipped, 0 failed
8 cases: 8 passed, 0 skipped, 0 failed
4 cases: 4 passed, 0 skipped, 0 failed
8 cases: 8 passed, 0 skipped, 0 failed
4 cases: 4 passed, 0 skipped, 0 failed
8 cases: 8 passed, 0 skipped, 0 failed
4 cases: 4 passed, 0 skipped, 0 failed
DONE
SCOPING FOR FAILURES IN CBLAS TESTS:
```

```

fgrep -e fault -e FAULT -e error -e ERROR -e fail -e FAIL \
    interfaces/blas/C/testing/Linux_ATHLON/sanity.out | \
    fgrep -v PASSED
make[1]: [sanity_test] Error 1 (ignored)
DONE
SCOPING FOR FAILURES IN F77BLAS TESTS:
fgrep -e fault -e FAULT -e error -e ERROR -e fail -e FAIL \
    interfaces/blas/F77/testing/Linux_ATHLON/sanity.out | \
    fgrep -v PASSED
make[1]: [sanity_test] Error 1 (ignored)
DONE

```

So, in the LAPACK testers we see no failures (all tests show 0 failed), and we have no output from the BLAS testers, which is what we want. Notice the lines like:

```
make[1]: [sanity_test] Error 1 (ignored)
```

This is due to fgrep’s behavior, and does not indicate an error. If fgrep does not find any pattern matches, it returns a 1, 0 on match. Therefore, since we are grepping for error, getting an “error condition” of 1 is what we hope for.

7.3 Finding the context of the error

If the sanity test output shows errors, the next step is to track down where they are coming from. You can see in the output the files that are being searched for errors. They are:

```

bin/sanity.out
interfaces/blas/C/testing/sanity.out
interfaces/blas/F77/testing/sanity.out

```

The threaded sanity test uses the same filenames with **pt** prefixed.

The first thing to notice is which of these tests are showing errors. The testers in bin are higher level than those in the interfaces directories, so if you get errors in both, track down and fix the interface errors first, as they may be causing the lapack errors. If both C and F77 BLAS interfaces are showing errors, I always scope and fix the Fortran77 stuff first, since Fortran is simpler (no RowMajor case to handle). Only if an error only shows up in C testing do I scope that output instead of the Fortran77.

The grepped error message probably gives you no idea what actually went wrong (it may show something as simple as:

```
FAIL
```

for instance), so you must go look at the **sanity.out** in question. For instance, you might need to scope **interfaces/blas/F77/testing/sanity.out**. You do a search for whatever alerted you to the problem (eg., FAIL), and you see by the surrounding context what tester failed.

7.4 Tracking down an error in the BLAS interface testers

The BLAS testers are split by BLAS Level (1, 2 or 3) and precision/type (s,d,c,z). The basic names of the tester executables are

```
x<pre>blat<lvl>
x<pre>cblat<lvl>
```

for Fortran77 and C, respectively. The Level 1 testers (`x[s,d,c,z]blat1`) test certain fixed cases, and thus take no input file. So if the error is in them, you simply run the executable with no args in order to reproduce the failure.

The Level 2 and 3 testers allow a user to specify what tests should be run, via an input file. The standard input files that ATLAS runs with are:

```
<pre>blat<lvl>.dat
c_<pre>blat<lvl>.dat
```

respectively. The format of these input files is pretty self explanatory, and more explanation can be found at:

www.netlib.org/blas/faq.html

To run the tester with these files, you redirect them into the tester. For instance, to run the double precision Level 2 tester with the default input file, you'd issue:

```
./xdblbat2 < ~/ATLAS/interfaces/blas/F77/testing/dblat2.dat
```

You should be aware that only the first error report in a run is accurate: one error can cause a cascade of spurious error reports, all of which may go away by fixing the first reported problem. So, it is important to find and fix the errors in sequence.

I usually copy the input file in question to a new file that I can hack on (for instance, if the error was in the double precision Level 2, I might issue:

```
cp ~/ATLAS/interfaces/blas/F77/testing/dblat2.dat bad.dat
```

I then repeatedly run the routine and simplify the input file until I have found the smallest, simplest input that displays the error.

The next step is to rule out tester error. The way I usually do this is to demonstrate that the error goes away by linking to the Fortran77 reference BLAS rather than ATLAS (you can only do this for errors in the F77 interface, obviously). I usually just do it by hand, i.e., for the same example again, I'd do:

```
f77 -o xtst dblat2.o /home/rwhaley/lib/libfblas.a
```

If the ATLAS-linked code has the error, and this one does not, it is a strong indication that the error is in ATLAS. If the F77 BLAS are shown to be in error, it is usually a compiler error, and can be fixed by turning down (or off) the optimization used to compile the tester.

Now you should have confirmed the tester is working properly, and that the error is in a specific routine (let us say DNRM2 as an example). As a quick proof that DNRM2 is indeed the problem, you can link explicitly to the F77 version of DNRM2, and to ATLAS for everything else (see Section ?? for hints on how to do this). If this still shows the error, you are confident that ATLAS's DNRM2 is indeed causing the problem, and you should either track it down, or report it (depending on your level of expertise).

7.5 Tracking down an error in the bin/ testers

The sanity tests only run the LAPACK testers in this directory. The LAPACK routines depend on the BLAS, so ignore errors in lapack testers until all the BLAS pass with no error. If you have errors in LAPACK but the BLAS pass all tests, then you have to hunt for the error in the LAPACK routines.

First, rule out that it's not a problem in the BLAS that is just not showing up in the BLAS testing. Get yourself a reference BLAS library, as explained in Section ?? . Then, set your `Make.inc`'s `BLASlib` macro to point to the created reference BLAS library. Then, you need to compile a library that uses ATLAS's lapack routines, but the reference BLAS. This can be done by compiling the same executable name with `_sys` suffixed. For instance, if you were running the LU tester, `xdlutst`, you would say `make xdlutst_sys`, and then run this executable with the same input.

If the error goes away, then the error is really in the ATLAS BLAS somewhere. I then usually look at the LAPACK routine and tester in question to find out what its BLAS dependencies are, and manually link in the reference BLAS object files until I find the exact BLAS causing the problem. Usually once you know what routine causes the prob, you can reproduce the error with the BLAS tester (i.e. you need a IDAMAX call with `N=12`, `incX=82`).

If the error still persists using ATLAS's LAPACK and the Fortran77 BLAS, the next trick is to do LAPACK just like the BLAS: download and compile the F77 LAPACK from netlib (www.netlib.org/lapack/lapack.tgz). You then set your `Make.inc`'s `FLAPACKlib` to point to your Fortran77 lapack library. You then suffix the base executable name with `F_sys` (eg., for LU again, you would do `make xdlutstF_sys`), and you will get a tester linked against the Fortran77 BLAS *and* LAPACK. If this also shows to be in error, there is an error in the tester, or in the compiler. Try turning down compiler optimization to rule in or out compiler errors.

8 Antoine's testing scripts

Before a stable release, we always do as much testing as possible. The 900 pound gorilla of testers is Antoine's tester scripts. This tester can run as long as several days, and does a great number of both fixed and random tests, and if it completes with no errors, you have a pretty good idea that the code is fairly solid. Even the casual user ought to run the sanity testing as a matter of course, and that should always be ran and passed first. Also, much of the methodology for understanding output, tracking down problems, etc, is the same for this tester and the sanity test, so read those sections first for tips I will not bother to repeat here.

You can run this full tester with this simple command given in your `BLDdir`:

```
make full_test
```

This can run for as long as several days (on modern machines, it usually takes more like 4-8 hours). If your connection to the machine is not stable, it can die before the test is complete, so it makes sense to `nohup` this command and redirect the output. I usually do (using `tcsh`):

```
nohup make full_test >& make.out &
```


Once you are done, you can auto-search for error messages with:

```
make scope full_test
```

8.1 Finding errors

If the scope finds problems, you will need to investigate the tester log files for more info. These files will be found in the directory:

```
BLDdir/bin/AtlasTest/<ARCHNAME>
```

the `.out` file in that directory should be a complete text capture of the full testing log. More detailed output can be found in subdirectories under here.

Once you have the error, you need to repeat it. You can try running the exact case, but sometimes that won't do it (for instance, you have a memory error that requires you to run many cases); you then need to find a small run that does demonstrate the error.

You should then apply the normal tricks (linking to F77 BLAS instead of sys blas, having the tester call the f77 blas twice, etc) to ensure the error really is in ATLAS, before tracking the error to its source.

8.2 Tracking down errors in the bin/ testers

There are two types of bin/ testers: lapack and blas. The BLAS testers have executable names of the form

```
x<pre>l<level>blastst
```

The BLAS testers test ATLAS against a known-good implementation, so the first thing to do is make sure the error is in ATLAS, and not the known-good implementation. To do this, change your BLASlib back to the default F77 reference BLAS, and if necessary compile them without optimization. and then relink and rerun the test in question. If the error goes away, you have found an error in your known-good library, not ATLAS. If it stays, you have found an error in ATLAS, and you should track it down or report it. See Section 7.5 for information on tracking problems in the LAPACK testers.

9 Finding a good NB for GEMM

One of the things I do most frequently with user-submitted kernels is reduce the blocking factor that the user has chosen. I often choose smaller NB than the best for asymptotic GEMM performance, and even more often choose one that does not yield the best performance in the kernel timer. To understand why, you must understand the following points, explained in turn below:

1. Better kernel timing (eg. `make ummcase` in your `<OBJdir>/tune/blas/gemm/` directory) does not always yield better total GEMM performance
2. Large NB means significantly more time in cleanup code
3. Large NB means significantly more time in unblocked application code

9.1 Better kernel timing does not always yield faster GEMM

The kernel timer (invoked by one of the `make mmcase` variants available in `<OBJdir>/tune/blas/gemm/`) tries to mimic the way ATLAS calls the kernel. However, it does not do everything the same way. First, there is no cleanup, so it is always calling the kernel only. More importantly, `CacheEdge` has not yet been determined, so no Level 2 Cache blocking is being performed. Therefore, it may sometimes look like you are better off to block the kernel for the L2 when using these kernel timers, when in fact, if you instead block for the Level 1 cache, `CacheEdge` will then further speed things up later, and thus the smaller NB achieves better GEMM performance, even when it runs slower in the kernel timer.

For machines with very large L1 caches, often several blocking factors that fit into L1 have roughly the same performance. In such a case, it is very likely that you want to choose the smallest achieving that rough performance, as it will allow more blocks to fit into the L2 blocking to be done later.

If a kernel appears to get much better performance with a large NB, the best idea is to build a full GEMM using both the best-performing small NB, and the best performing large NB, and seeing what the gap truly is. Very often, the small kernel will actually be better even asymptotically, and if it is not, it will often be so much better for smaller problems that it makes sense to use it anyway.

Even beyond these explanations, it is sometimes the case that the kernel timer predicts good performance that is not realized when the full GEMM is built. This is usually due to inadequate cache flushing, leading to overprediction of performance because things are retained more in the cache than they are in practice. Therefore, I usually pump up the flushing mechanism (set `L2SIZE` of your `Make.inc` to ridiculously large levels). No matter what, actual full GEMM performance is the final arbiter. If it is not as high as predicted by the kernel timer, it may be worthwhile to see if other, smaller NB, cases achieve the same full-gemm performance.

9.2 Large NB means more time in cleanup

One bad news about choosing a large NB is that applications will spend more of their time in cleanup. Let us say you choose a block factor of 120. In this case, many applications will never even call your optimized kernel, but spend all their time in GEMM cleanup. Some applications are statically blocked, and if their NB is smaller than yours, they can spend their entire time in cleanup even for large problems.

Therefore, if you must choose a large NB in order to get adequate GEMM performance, you must pay an unusual amount of attention to cleanup optimization. However, as the next section will discuss, even if cleanup ran at the same speed as your best kernel, this will yield poor performance for many codes.

9.3 Large NB means more time in unblocked application code

Probably the worst thing about choosing a large NB is that many applications use Level 1 and 2 BLAS in order to do the unblocked part of the computation. These BLAS are usually at least an order of magnitude slower than GEMM. Therefore, as you increase NB, for applications with unblocked portions, you increase the proportion of time spent in this order-of-magnitude slower code. Therefore, even with perfect cleanup, a large NB may result

in an application running **at less than half speed**, even though GEMM performance is quite good.

To get an idea of this, simply scope the factorizations provided by LAPACK. These applications are statically blocked, so that the column factorizations (eg., DGETF2 for LU) are used until N_B is reached. If ILAENV returns a blocking factor smaller than your GEMM, the applications will stay in cleanup even for large problems. Even worse, some applications (eg., QR) require workspace proportional to N_B , and since dynamic memory is not used, it is possible even if you hack ILAENV to use the correct blocking factor, they will be forced to a smaller one.

9.4 Finding a good N_B

I will call the first level of cache accessed by the floating point unit the Level 1 cache, regardless of whether it is the first level of cache of the machine (there are a number of machines, such as the P4 (prescott) and Itanium where the FPU skips the Level 1 cache). Let N_e be the number of elements of the data type of interest in this cache. If this cache is write-through, then a rough guess for a good upper bound is $N_B \leq \sqrt{N_e}$. If the cache is not write-through, this is still the upper bound, but many larger caches often benefit from using a smaller N_B , one roughly $N_B < \frac{\sqrt{N_e}}{3}$. We can describe this more exactly, but these bounds are easy to compute during tuning.

You should not choose an N_B that is a power of 2, as this could occasionally cause nasty cache conflicts. There's often a small advantage to choosing N_B that are a multiple of cache line size; this can sometimes be critical, depending on the arch.

So, the basic idea is to start looking at N_B given by the above two computations, and then try a little smaller and larger using the kernel timer. If you get two that tie for out-of-cache performance, always take the smaller. If best performance is achieved with very large N_B (say $N_B \geq 80$), then always confirm that it yields better GEMM performance than a smaller N_B , and that application performance is not severely impacted, particularly for smaller problems.

The way I usually time application performance is to time ATLAS's LU. This actually gives you a very rosy picture of how a large block factor will effect performance, in that it uses recursion rather than statically blocking. This means that ATLAS's LU does not have any unblocked code, and thus doesn't slow down the way LAPACK's LU will for large N_B . However, if even this code shows performance loss for smaller sizes, you know your cleanup needs to get a *lot* better, or you need to reduce N_B , even if it results in a slight reduction in GEMM performance. If you want to get a better idea of how most applications will perform, time one of LAPACK's factorizations instead.

Under no circumstances should you choose a blocking factor much larger than 120. I confine the ATLAS search to a maximal size of 80 for the above reasons, but occasionally go a little higher for machines without effective L1 caches. However, this can absolutely kill application performance. Further, it is never a good idea to completely fill an Level 2 cache with your block. It may look good in GEMM, but it will die in any application, both for the reasons above, and the following: The L2 cache is shared instruction/data. Filling it with data will often lead to instruction loading/flushing cycle when a larger application is calling. Remember that GEMM is of interest because of all the applications that are built from it, not when used in isolation.

If a NB larger than 60 only gives you a few percent, always choose a smaller one; only go above 80 for significant advantage, and essentially don't go above 120 unless absolutely necessary, and then you can expect slowdown in many applications, even once you have fully optimized all cleanup cases.

10 Understanding the ATLAS configure system

NOTE: this information was out of date before it was finished, so this discussion should be viewed as an introduction only.

For ATLAS 3.7.12, ATLAS's configure routine was completely rewritten for greater modularity. The total amount of code probably increased, but the amount that must be examined at any time should be very much smaller.

In the new system, the topmost unit is `ATLAS/configure` which is a BFI shell script which allows ATLAS's `config.c` to be invoked in a way very similar to `gnu configure`. This shell script gathers some info and fills in a `Makfile` which is then used to build `xconfig` from `ATLAS/CONFIG/src/config.c`. `config.c` is a driver program that first calls various probes to determine any information not overridden by user flags, and then calls `xspew` to create a full `Make.inc` for the target architecture. `xspew` is built from the file `ATLAS/CONFIG/src/SpewMakeInc.c`.

The idea is to change ATLAS's install so it consists of the following commands:

1. `/path/to/ATLAS/configure` : Create `Make.inc` and build subdirs in the present directory (ATLAS no longer requires building in arch-spec directories under the source tree)
2. `make build` : Build ATLAS
3. `make check` : run sanity tests
4. `make time` : run simple benchmarks, compare observed vs. expected performance, and issue warning if too low
 - Keep record of arch default installs with performance as % of clock rate (not peak!)
 - If arch def used wt bad compiler, this will detect performance difference, and user should be warned
 - If no arch defaults (so expected performance unknown), still create standard benchmark output file for submission to ATLAS
5. `make install` : copy libraries and include files to user-specified directories

Note that steps 2-5 are all handled by the build framework, so here we are only talking about step 1 (the configure framework).

10.1 Probe Overview

From ATLAS 3.7.12 on, ATLAS's config routine was rewritten for greater modularity, with each config probe having its own driver and so on. For this discussion, we will refer to the machine doing the cross-compilation as the *frontend* (abbreviated as FE), and the machine which ATLAS is being tuned for the *backend* (abbreviated as BE). Note that if you are not doing cross-compilation (the majority of the time) the front-end and back-end are the same machine.

Every type of probe has a frontend driver (occasionally, config may directly call the backend driver, if there is only one) which will itself call multiple backend drivers. For

instance, the probe to compute the architecture runs on the frontend, and calls different backend drivers depending on the assembly dialect and operating system of the backend. The files for the frontend drivers are located in `ATLAS/CONFIG/src`, and the backend files are in `ATLAS/CONFIG/src/backend`, with all include files in `ATLAS/CONFIG/include`. All frontend probes use the file `atlconf_misc.c` (prototyped in `atlconf_misc.h`), which handles things like file I/O, issuing shell commands, etc. Some of the current probes used by config are:

1. OS Probe

Purpose: Discover the Operating System being used

Inputs: None

Outputs: Enumerated type of OS

FE files: `probe_OS.c`

BE files: None (`uname` on BE)

2. Assembly dialect probe

Purpose: Discover what ATLAS assembly dialect works

Inputs: OS enum (gives subdialect of assembler)

Outputs: Enum of assembly dialect

FE files: `probe_asm.c`,

BE files: `probe_this_asm.c` – [`probe_gas_parisc.S`, `probe_gas_ppc.S`, `probe_gas_sparc.S`, `probe_gas_x8632.S`, `probe_gas_x8664.S`]

3. Vector ISA extension probe – assembly

Purpose: Discover which of supported vector ISA extensions work

Inputs: enums for OS and assembly dialect

Outputs: `iflag = ((1<<ISA0) | (1<<ISA1) | ... | (1<<ISAn))`

FE files: `probe_vec.c`

BE files: `probe_svec.c` – [`probe_AltiVec.S`, `probe_SSE.S`], `probe_dvec.c` – [`probe_SSE2.S`], `probe_dSSE3.c` – [`probe_SSE3.S`]

4. Architecture probe

Purpose: Discover target architecture/machine

Inputs: OS and assembly enums [force 64/32 bit usage]

Outputs: enum of arch

FE files: `archinfo.c`

BE files: `archinfo_x86.c`, `archinfo_linux.c`, `archinfo_freebsd.c`, `archinfo_aix.c`, `archinfo_irix.c`, `archinfo_sunos.c`

Notes: See Section 10.2 for more details.

5. 64-bit probe

Purpose: Discover if arch supports 64-bit pointers

Inputs: OS, arch [user choice]

Outputs: 32 / 64

files: Config directly calls archinfo

Notes: : New policy: config assumes whatever compiler gives you w/o -m32 -m64, and user must throw special flag to append these to the line.

6. Compiler probe

Purpose: Find good compilers

Inputs: OS, arch [,suggested compilers]

Outputs: The following:

- (a) F2CNAME, F2CINT, F2CSTRING enums
- (b) Compilers and flags

FE files : probe_comp.c probe_f2c.c probe_ccomp.c

BE files : f2cname[F,C] . [f,c], f2cint[F,C] . [f,c], f2cstr[F,C] . [f,c], ccomp
interaction not yet done

Note: This is complex, see Section 10.4 for details.

7. Arch defaults probe

Purpose: Discover arch defaults

Inputs: OS, arch, compilers

Outputs: Whether to use arch defs (INSTFLAGS in Make.inc)

files: ARCHS/Makefile

invoke: Arch default setup is instigated by atlas_install.c.

notes: May want to have it autobenchmark kernel, test against table of expected perf, to see if arch def are OK wt this compiler version.

10.2 Architectural Probes

We use the `archinfo_xxx` probes to discover the following architectural information:

- 'n': number of cpus
- 'c': number of cache levels
- 'C' #: size in KB of cache level #
- 'v': verbose (prints strings as well as ints)
- 'm': clock rate in Mhz
- 'a': ATLAS architecture classification
- 'b': support for 64 / 32 bits
- 't': Is cpu throttling currently on (-1: no, 0: don't know, 1: yes)

If a given probe cannot find that particular item, it is returned as 0.

The frontend wrapper script `archinfo.c` calls these guys according to OS, and tries to get all flags filled in with union of functionality of `archinfo_x86` and `archinfo_<OS>`.

10.3 Notes on configure

New policies:

- Any 64-bit arch defaults to building whatever the compiler natively does w/o flags unless overridden
- L2SIZE always set to 4MB unless overridden
- Timer defaults to standard wall/cpu unless overridden
 - Exception is solaris, where we default to solaris hr timers
 - If Mhz is passed in, use cycle-accurate wall-timer on x86

10.4 Compiler Setup and Handling in ATLAS Config

This is complicated as hell. Potentially, each architecture/OS combo has unique compiler and flags for each supported compiler (more below), and the user can override any/all of these. I'm changing the number of supported compilers for greater flexibility (all compilers except XCC are for the backend system, in the case of cross-compilation). These are:

ICC : compiles all C interface routines. Since it is not used for any kernel compilation the performance impact of this compiler should be minimal.

SMC : used to compile ATLAS single precision matmul kernels

DMC : used to compile ATLAS double precision matmul kernels

SKC : used to compile all non-interface, non-gemm-kernel single precision ATLAS routines

DKC : used to compile all non-interface, non-gemm-kernel double precision ATLAS routines

XCC : used to compile all front-end codes

GCC : used to assemble assembly files and to compile kernels that call for gcc

F77 : Valid fixed-format Fortran77 compiler that compiles ATLAS's F77 interface routines. This should match the Fortran77 the user is using. This compiler's performance does not affect ATLAS's performance, and so may be anything.

Here's my present design:

1. **Compiler defaults:** are read in from `atlcomp.txt`, which allows the user specify default compiler/flags, as well as specific ones for particular architectures, and multiple compilers for a given arch.
 - `ATLAS/SRC/probe_comp.c`
2. **Executable search:** takes name of executable (in this case a compiler name), and finds the path to it. Skipped if the user provides the path as part of the compiler.

3. **C compiler interaction probe:** separate probe that takes two or more C compilers and their flags as arguments, and makes sure they are able to call each other w/o problems.

- *unimplemented:* presently assumed to work

4. **F77/C calling convention probe:** as in present config, but as an independent probe.

- *done:* front-end is `probe_f2c.c`

10.5 Notes on cross-compilation:

For this discussion, we will refer to the machine doing the cross-compilation as the *frontend* (abbreviated as FE), and the machine which ATLAS is being tuned for the *backend* (abbreviated as BE). If you wish to use ATLAS in cross-compilation mode, you will usually throw the configure flag `--rtarg=<mach>`, where `<mach>` is a valid machine name or IP address of the BE. For this to work, you should be logged into the FE, and you must be able to ssh to the BE from there without prompting for a password, and you must be installing ATLAS in a shared filesystem (it should be possible to remove the requirement that systems have a shared filesystem, but so far we haven't seen the necessity).

The `configure` script will always produce a file `BLDdir/ATLrun.sh`, which will always be invoked when executing programs that are meant for the BE. If cross-compilation is not being used, this shell script will simply execute the routine locally, otherwise `configure` will generate a slightly more complicated script that spawns the executable using `ssh` to the BE. Currently, `config.c`'s routines have arguments that pass around the BE information, but this is no longer needed, and so is never used (this code is from earlier support for cross-compilation that has been replaced by the `ATLrun.sh` trick, and will be removed the next time `config.c` is refactored).

10.6 Notes on native Windows compilation:

Native Windows compilers like MSVC++ or Intel's `icc` do not take unix-style flags or paths. This is a problem, because ATLAS uses a huge amount of Makefiles that assume unix-style flags and paths. We partially solve this problem by requiring cygwin to be installed on Windows, which provides a 32-bit `gcc`, `make`, and unix-style filesystem paths. However, the user may wish to build ATLAS with a native windows compiler, which, not knowing about cygwin, will be confused by the passed-in flags and paths. Therefore, on Windows, native-windows compilers are not called directly by ATLAS. Instead, during the configure process ATLAS compiles small C programs that wrap the compilers and translate from unix-style to windows-style components.

The simplest of these is when we need to use the MinGW compilers for building ATLAS. The MinGW compilers are simply the GNU compilers that have been made native to Windows (they use the Windows `libc/threads/etc.`), so they no longer depend on the cygwin environment at all. This means that MinGW compilers produce executables that don't require cygwin, and thus ATLAS compiled with MinGW can run on Windows system without cygwin installed, and don't cause any more dependencies than using a native compiler.